

Finetuning Qwen2.5-7B-Instruct model using LoRA and QLoRA

```
# pip installs
```

```
!pip install -q --upgrade torch==2.5.1+cu124 torchvision==0.20.1+cu124 torchaudio==2.5.1+cu124 --index-u
!pip install -q requests bitsandbytes==0.46.0 transformers==4.48.3 accelerate==1.3.0
!pip install -q datasets requests peft
```

```
===== 908.2/908.2 MB 1.2 MB/s eta 0:00:00
===== 7.3/7.3 MB 35.1 MB/s eta 0:00:00
===== 3.4/3.4 MB 43.5 MB/s eta 0:00:00
===== 24.6/24.6 MB 71.1 MB/s eta 0:00:00
===== 883.7/883.7 kB 40.4 MB/s eta 0:00:00
===== 13.8/13.8 MB 73.0 MB/s eta 0:00:00
===== 664.8/664.8 MB ? eta 0:00:00
===== 363.4/363.4 MB 1.2 MB/s eta 0:00:00
===== 211.5/211.5 MB 6.7 MB/s eta 0:00:00
===== 56.3/56.3 MB 11.7 MB/s eta 0:00:00
===== 127.9/127.9 MB 8.9 MB/s eta 0:00:00
===== 207.5/207.5 MB 1.6 MB/s eta 0:00:00
===== 188.7/188.7 MB 6.8 MB/s eta 0:00:00
===== 99.1/99.1 kB 9.2 MB/s eta 0:00:00
===== 21.1/21.1 MB 112.0 MB/s eta 0:00:00
===== 209.6/209.6 MB ? eta 0:00:00
===== 6.2/6.2 MB 81.4 MB/s eta 0:00:00
===== 44.4/44.4 kB 3.2 MB/s eta 0:00:00
===== 67.0/67.0 MB 12.9 MB/s eta 0:00:00
===== 9.7/9.7 MB 138.4 MB/s eta 0:00:00
===== 336.6/336.6 kB 29.6 MB/s eta 0:00:00
===== 3.1/3.1 MB 98.2 MB/s eta 0:00:00
```

```
# imports
```

```
import os
import re
import math
from tqdm import tqdm
from google.colab import userdata
from huggingface_hub import login
import torch
import transformers
from transformers import AutoModelForCausalLM, AutoTokenizer, BitsAndBytesConfig, TrainingArguments, se
from peft import LoraConfig, PeftModel
from datetime import datetime
```

```
# Constants
```

```
BASE_MODEL = "Qwen/Qwen2.5-7B-Instruct"
```

```
# Hyperparameters for QLoRA Fine-Tuning
```

```
LORA_R = 32
```

```
LORA_ALPHA = 64
```

```
TARGET_MODULES = ["q_proj", "v_proj", "k_proj", "o_proj"]
```

Log in to HuggingFace

```
# Log in to HuggingFace
```

```
hf_token = userdata.get('HF_TOKEN')
login(hf_token, add_to_git_credential=True)
```



Analyzing the base model

```
# Load the Base Model without quantization
```

```
base_model = AutoModelForCausalLM.from_pretrained(BASE_MODEL, device_map="auto")
```

```
config.json: 100% 663/663 [00:00<00:00, 77.7kB/s]
```

```
model.safetensors.index.json: 27.8k/? [00:00<00:00, 2.17MB/s]
```

```
Downloading shards: 100% 4/4 [04:32<00:00, 67.03s/it]
```

```
model-00001-of-00004.safetensors: 100% 3.95G/3.95G [01:03<00:00, 54.8MB/s]
```

```
model-00002-of-00004.safetensors: 100% 3.86G/3.86G [01:22<00:00, 75.0MB/s]
```

```
model-00003-of-00004.safetensors: 100% 3.86G/3.86G [01:01<00:00, 120MB/s]
```

```
model-00004-of-00004.safetensors: 100% 3.56G/3.56G [01:03<00:00, 60.8MB/s]
```

```
Loading checkpoint shards: 100% 4/4 [00:55<00:00, 18.24s/it]
```

```
generation_config.json: 100% 243/243 [00:00<00:00, 26.4kB/s]
```

```
WARNING:accelerate.big_modeling:Some parameters are on the meta device because they were offloaded to th
```

Double-click (or enter) to edit

```
print(f"Memory footprint: {base_model.get_memory_footprint() / 1e9:,.1f} GB")
```

```
Memory footprint: 30.5 GB
```

```
base_model
```

```
Qwen2ForCausalLM(
  (model): Qwen2Model(
    (embed_tokens): Embedding(152064, 3584)
    (layers): ModuleList(
      (0-27): 28 x Qwen2DecoderLayer(
        (self_attn): Qwen2Attention(
          (q_proj): Linear(in_features=3584, out_features=3584, bias=True)
          (k_proj): Linear(in_features=3584, out_features=512, bias=True)
          (v_proj): Linear(in_features=3584, out_features=512, bias=True)
          (o_proj): Linear(in_features=3584, out_features=3584, bias=False)
        )
        (mlp): Qwen2MLP(
          (gate_proj): Linear(in_features=3584, out_features=18944, bias=False)
          (up_proj): Linear(in_features=3584, out_features=18944, bias=False)
          (down_proj): Linear(in_features=18944, out_features=3584, bias=False)
          (act_fn): SiLU()
        )
        (input_layernorm): Qwen2RMSNorm((3584,), eps=1e-06)
        (post_attention_layernorm): Qwen2RMSNorm((3584,), eps=1e-06)
      )
    )
    (norm): Qwen2RMSNorm((3584,), eps=1e-06)
    (rotary_emb): Qwen2RotaryEmbedding()
  )
  (lm_head): Linear(in_features=3584, out_features=152064, bias=False)
)
```

Python 3 Google Compute Engine backend (GPU)

Showing resources from 23:09 to 23:23

System RAM
10.5 / 12.7 GBGPU RAM
10.9 / 15.0 GBDisk
71.0 / 112.6 GB

Quantizing with 8 bits

```
# Load the Base Model using 8 bit

quant_config = BitsAndBytesConfig(load_in_8bit=True)

base_model = AutoModelForCausalLM.from_pretrained(
    BASE_MODEL,
    quantization_config=quant_config,
    device_map="auto",
)
```

Loading checkpoint shards: 100%

4/4 [01:12<00:00, 17.54s/it]

Double-click (or enter) to edit

```
print(f"Memory footprint: {base_model.get_memory_footprint() / 1e9:.1f} GB")
```

Memory footprint: 8.7 GB

base_model

```
Qwen2ForCausalLM(
  (model): Qwen2Model(
    (embed_tokens): Embedding(152064, 3584)
    (layers): ModuleList(
      (0-27): 28 x Qwen2DecoderLayer(
        (self_attn): Qwen2Attention(
          (q_proj): Linear8bitLt(in_features=3584, out_features=3584, bias=True)
          (k_proj): Linear8bitLt(in_features=3584, out_features=512, bias=True)
          (v_proj): Linear8bitLt(in_features=3584, out_features=512, bias=True)
          (o_proj): Linear8bitLt(in_features=3584, out_features=3584, bias=False)
        )
        (mlp): Qwen2MLP(
          (gate_proj): Linear8bitLt(in_features=3584, out_features=18944, bias=False)
          (up_proj): Linear8bitLt(in_features=3584, out_features=18944, bias=False)
          (down_proj): Linear8bitLt(in_features=18944, out_features=3584, bias=False)
          (act_fn): SiLU()
        )
        (input_layernorm): Qwen2RMSNorm((3584,), eps=1e-06)
        (post_attention_layernorm): Qwen2RMSNorm((3584,), eps=1e-06)
      )
    )
    (norm): Qwen2RMSNorm((3584,), eps=1e-06)
    (rotary_emb): Qwen2RotaryEmbedding()
  )
  (lm_head): Linear(in_features=3584, out_features=152064, bias=False)
)
```

Python 3 Google Compute Engine backend (GPU)

Showing resources from 23:09 to 23:30

System RAM
3.5 / 12.7 GBGPU RAM
8.3 / 15.0 GBDisk
71.0 / 112.6 GB

Double-click (or enter) to edit

Double Quantization with 4 bits

Load the Tokenizer and the Base Model using 4 bit

```
quant_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_use_double_quant=True,
    bnb_4bit_compute_dtype=torch.bfloat16,
    bnb_4bit_quant_type="nf4")
```

```
base_model = AutoModelForCausalLM.from_pretrained(
    BASE_MODEL,
    quantization_config=quant_config,
    device_map="auto",
)
```

Loading checkpoint shards: 100%

4/4 [01:15<00:00, 18.41s/it]

```
print(f"Memory footprint: {base_model.get_memory_footprint() / 1e9:,.2f} GB")
```

Memory footprint: 5.44 GB

base_model

```
Qwen2ForCausalLM(
  (model): Qwen2Model(
    (embed_tokens): Embedding(152064, 3584)
    (layers): ModuleList(
      (0-27): 28 x Qwen2DecoderLayer(
        (self_attn): Qwen2Attention(
          (q_proj): Linear4bit(in_features=3584, out_features=3584, bias=True)
          (k_proj): Linear4bit(in_features=3584, out_features=512, bias=True)
          (v_proj): Linear4bit(in_features=3584, out_features=512, bias=True)
          (o_proj): Linear4bit(in_features=3584, out_features=3584, bias=False)
        )
        (mlp): Qwen2MLP(
          (gate_proj): Linear4bit(in_features=3584, out_features=18944, bias=False)
          (up_proj): Linear4bit(in_features=3584, out_features=18944, bias=False)
          (down_proj): Linear4bit(in_features=18944, out_features=3584, bias=False)
          (act_fn): SiLU()
        )
        (input_layernorm): Qwen2RMSNorm((3584,), eps=1e-06)
        (post_attention_layernorm): Qwen2RMSNorm((3584,), eps=1e-06)
      )
    )
    (norm): Qwen2RMSNorm((3584,), eps=1e-06)
    (rotary_emb): Qwen2RotaryEmbedding()
  )
)
```

```
(lm_head): Linear(in_features=3584, out_features=152064, bias=False)
)
```

Python 3 Google Compute Engine backend (GPU)

Showing resources from 23:09 to 23:34

System RAM
3.6 / 12.7 GB



GPU RAM
5.4 / 15.0 GB



Disk
71.0 / 112.6 GB

